



Constraint Satisfaction Problems (CSPs)


Russell and Norvig Chapter 5




CSP example: map coloring




Given a map of Australia, color it using three colors such that no neighboring territories have the same color.



CSP example: map coloring



- Solutions are assignments satisfying all constraints, e.g.:
 $\{WA=red, NT=green, Q=red, NSW=green, V=red, SA=blue, T=green\}$

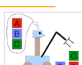


Constraint satisfaction problems

- A CSP is composed of:
 - A set of variables X_1, X_2, \dots, X_n with domains (possible values) D_1, D_2, \dots, D_n
 - A set of constraints C_1, C_2, \dots, C_m
 - Each constraint C_i limits the values that a subset of variables can take, e.g., $V_1 \neq V_2$


In our example:

- Variables: WA, NT, Q, NSW, V, SA, T
- Domains: $D_i = \{red, green, blue\}$
- Constraints: adjacent regions must have different colors.
 - E.g. $WA \neq NT$ (if the language allows this) or
 - $\{WA, NT\}$ in $\{(red, green), (red, blue), (green, red), (green, blue), (blue, red), (blue, green)\}$




Constraint satisfaction problems

- A **state** is defined by an assignment of values to some or all variables.
- Consistent assignment**: assignment that does not violate the constraints.
- Complete assignment**: every variable is mentioned.
- Goal: a complete, legal assignment.



$\{WA=red, NT=green, Q=red, NSW=green, V=red, SA=blue, T=green\}$



Constraint satisfaction problems

- Simple example of a **formal representation language**
- CSP benefits
 - Standard representation language
 - Generic goal and successor functions
 - Useful **general-purpose** algorithms with more power than standard search algorithms, including generic heuristics
- Applications:
 - Time table problems (exam/teaching schedules)
 - Assignment problems (who teaches what)

Varieties of CSPs

- Discrete variables
 - Finite domains of size $d \Rightarrow O(d^n)$ complete assignments.
 - The satisfiability problem: a Boolean CSP
 - Infinite domains (integers, strings, etc.)
 - E.g. job scheduling where variables are start/end days for each job.
 - Need a constraint language e.g. $StartJob_1 + 5 \leq StartJob_2$.
- Continuous variables
 - e.g. start/end times for Hubble Telescope observations.
 - Linear constraints solvable in poly time by linear programming methods (dealt with in the field of operations research).

September 28, 2009

7

Varieties of constraints


- Unary constraints involve a single variable.
 - e.g. $SA \neq green$
- Binary constraints involve pairs of variables.
 - e.g. $SA \neq WA$
- Higher-order constraints involve 3 or more variables.
- Preference (soft constraints) e.g. *red* is better than *green* often representable by a cost for each variable assignment; not considered here.

September 28, 2009

8

Constraint graph

- Binary CSP:** each constraint relates two variables
- Constraint graph:** nodes are variables, edges are constraints




September 28, 2009

9

Example: cryptarithmic puzzles

$$\begin{array}{r} T W O \\ + T W O \\ \hline F O U R \end{array}$$



Variables: $F T U W R O X_1 X_2 X_3$
Domains: $\{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$
Constraints
 $allDiff(F, T, U, W, R, O)$
 $O + O = R + 10 \cdot X_1$, etc.

September 28, 2009

10

CSP as a standard search problem

- Incremental formulation
 - Initial State:** the empty assignment $\{\}$.
 - Successor function:** Assign value to unassigned variable provided that there is not conflict.
 - Goal test:** the current assignment is complete.
- Same formulation for all CSPs !!!
- Solution is found at depth n (n variables).
 - What search method would you choose?

September 28, 2009

11

Backtracking search

- Observation: the order of assignment doesn't matter
 - \Rightarrow can consider assignment of a single variable at a time. Results in d^n leaves.
- Backtracking search: DFS for CSPs with single-variable assignments (backtracks when a variable has no value that can be assigned)
- The basic uninformed algorithm for CSP

September 28, 2009

12

Backtracking search

```

function BACKTRACKING-SEARCH(csp) return a solution or failure
  return RECURSIVE-BACKTRACKING( $\emptyset$ , csp)

function RECURSIVE-BACKTRACKING(assignment, csp) return a solution or failure
  if assignment is complete then return assignment
  var ← SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assignment, csp)
  for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
    if value is consistent with assignment according to
      CONSTRAINTS[csp] then
      add (var=value) to assignment
      result ← RECURSIVE-BACKTRACKING(assignment, csp)
      if result ≠ failure then return result
      remove (var=value) from assignment
  return failure
    
```

September 28, 2009

13

Backtracking example



September 28, 2009

14


Backtracking example



September 28, 2009

15


Backtracking example



September 28, 2009

16

Backtracking example



September 28, 2009

17

Improving backtracking efficiency

- General-purpose methods can give huge gains in speed:
 - Which variable should be assigned next?
 - In what order should its values be tried?
 - Can we detect inevitable failure early?

September 28, 2009

18

Most constrained variable




`var ← SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assignment, csp)`

Choose the variable with the fewest legal values
(most constrained variable)
a.k.a minimum remaining values (MRV) or "fail first" heuristic
□ What is the intuition behind this choice?

September 28, 2009

19

Most constraining variable



- Select the variable that is involved in the largest number of constraints on other unassigned variables.
- Also called the *degree* heuristic because that variable has the largest degree in the constraint graph.
- Often used as a tie breaker e.g. in conjunction with MRV.

September 28, 2009

20

Least constraining value




- Least constraining value heuristic: guides the choice of which value to assign next.
- Given a variable, choose the least constraining value:
 - the one that rules out the fewest values in the remaining variables
 - why?

September 28, 2009

21

Forward checking




- Can we detect inevitable failure early?
 - And avoid it later?
- *Forward checking*: keep track of remaining legal values for unassigned variables.
- Terminate search direction when a variable has no legal values.

September 28, 2009

22

Forward checking




- Assign {WA=red}
- Effects on other variables connected by constraints with WA
 - NT can no longer be red
 - SA can no longer be red

September 28, 2009

23

Forward checking




- Assign {Q=green}
- Effects on other variables connected by constraints with Q
 - NT can no longer be green
 - NSW can no longer be green
 - SA can no longer be green

September 28, 2009

24

Forward checking





- If V is assigned *blue*
- Effects on other variables connected by constraints with WA
 - SA is empty
 - NSW can no longer be *blue*
- FC has detected that partial assignment is *inconsistent* with the constraints and backtracking can occur.

September 28, 2009

25



Example: 4-Queens Problem



September 28, 2009

26



Example: 4-Queens Problem



September 28, 2009

27



Example: 4-Queens Problem



September 28, 2009

28



Example: 4-Queens Problem



September 28, 2009

29



Example: 4-Queens Problem



September 28, 2009

30



Example: 4-Queens Problem



September 28, 2009

31



Example: 4-Queens Problem



September 28, 2009

32



Example: 4-Queens Problem



September 28, 2009

33



Example: 4-Queens Problem



September 28, 2009

34



Example: 4-Queens Problem



September 28, 2009

35



Example: 4-Queens Problem



September 28, 2009

36

Example: 4-Queens Problem



September 28, 2009

37

Constraint propagation




- Solving CSPs with combination of heuristics plus forward checking is more efficient than either approach alone.
- FC does not provide early detection of all failures.
 - Once WA=red and Q=green: NT and SA cannot both be blue!
- Constraint propagation: propagate the implications of each constraint

September 28, 2009

38

Arc consistency




- $X \rightarrow Y$ is **consistent** iff for every value x of X there is some allowed y
- SA \rightarrow NSW is consistent since SA=blue and NSW=red is a consistent assignment.
- Arc – directed edge

September 28, 2009

39

Arc consistency




- $X \rightarrow Y$ is **consistent** iff for every value x of X there is some allowed y
- NSW \rightarrow SA is not consistent since for NSW=blue there is no consistent assignment to SA.
- Arc can be made consistent by removing blue from NSW

September 28, 2009

40

Arc consistency




- RECHECK neighbours !!
 - Remove red from V

September 28, 2009

41

Arc consistency



- Can be run as a preprocessing before the search or after each assignment.
 - Repeated until no inconsistency remains

September 28, 2009

42

Arc Consistency Algorithm

function AC-3(*csp*) **return** the CSP, possibly with reduced domains

inputs: *csp*, a binary csp with variables $\{X_1, X_2, \dots, X_n\}$

local variables: *queue*, a queue of arcs initially the arcs in *csp*

while *queue* is not empty **do**

$(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(\text{queue})$

if REMOVE-INCONSISTENT-VALUES(X_i, X_j) **then**

for each X_k **in** NEIGHBORS[X_j] **do**

 add (X_i, X_k) to *queue*

function REMOVE-INCONSISTENT-VALUES(X_i, X_j) **return** *true* iff we remove a value

removed \leftarrow *false*

for each x **in** DOMAIN[X_i] **do**

if no value y in DOMAIN[X_j] allows (x, y) to satisfy the constraints between X_i and X_j

then delete x from DOMAIN[X_i]; *removed* \leftarrow *true*

return *removed*

Time complexity: $O(n^2d^3)$

September 28, 2009

43

K-consistency

- Arc consistency does not detect all inconsistencies:
 - Partial assignment $\{WA=red, NSW=red\}$ is inconsistent.
- Stronger forms of propagation can be defined using the notion of k-consistency.
- A CSP is k-consistent if for any set of k-1 variables and for any consistent assignment to those variables, a consistent value can always be assigned to any kth variable.
- Not practical!

September 28, 2009

44


Local search for CSP

- Local search methods use a "complete" state representation, i.e., all variables assigned.
- To apply to CSPs
 - Allow states with unsatisfied constraints
 - operators **reassign** variable values
- Select a variable: random conflicted variable
- Select a value: *min-conflicts heuristic*
 - Value that violates the fewest constraints
 - Hill-climbing like algorithm with the objective function being the number of violated constraints
- Works surprisingly well in problem like n-Queens

September 28, 2009

45

Problem structure




- How can the problem structure help to find a solution quickly?
- Subproblem identification is important:
 - Coloring Tasmania and mainland are independent subproblems
 - Identifiable as connected components of constraint graph.
- Improves performance

September 28, 2009

46

Problem structure




- Suppose each problem has c variables out of a total of n .
- Worst case solution cost is $O(n/c \cdot d^c)$ instead of $O(d^n)$
- Suppose $n=80$, $c=20$, $d=2$
 - $2^{80} = 4$ billion years at 1 million nodes/sec.
 - $4 \cdot 2^{20} = .4$ second at 1 million nodes/sec

September 28, 2009

47

Tree-structured CSPs




- Theorem: if the constraint graph has no loops then CSP can be solved in $O(nd^2)$ time
- Compare with general CSP, where worst case is $O(d^n)$

September 28, 2009

48

Tree-structured CSPs




- Any tree-structured CSP can be solved in time linear in the number of variables.
 - Choose a variable as root, order variables from root to leaves such that every node's parent precedes it in the ordering. (label var from X_1 to X_n)
 - For j from n down to 2, apply REMOVE-INCONSISTENT-VALUES(Parent(X_j), X_j)
 - For j from 1 to n assign X_j consistently with Parent(X_j)

September 28, 2009

49

Nearly tree-structured CSPs




- Can more general constraint graphs be reduced to trees?
- Two approaches:
 - Remove certain nodes
 - Collapse certain nodes

September 28, 2009

50

Nearly tree-structured CSPs




- Idea: assign values to some variables so that the remaining variables form a tree.
- Assign $\{SA=x\} \leftarrow \text{cycle cutset}$
 - Remove any values from the other variables that are inconsistent.
 - The selected value for SA could be the wrong: have to try all of them

September 28, 2009

51

Nearly tree-structured CSPs



- This approach is effective if cycle cutset is small.
- Finding the smallest cycle cutset is NP-hard
 - Approximation algorithms exist
- This approach is called *cutset conditioning*.

September 28, 2009

52

Summary

- CSPs are a special kind of problem: states defined by values of a fixed set of variables, goal test defined by constraints on variable values
- Backtracking=depth-first search with one variable assigned per node
- Variable ordering and value selection heuristics help significantly
- Forward checking prevents assignments that lead to failure.
- Constraint propagation does additional work to constrain values and detect inconsistencies.
- Structure of CSP affects its complexity. Tree structured CSPs can be solved in linear time.

September 28, 2009

53